



# Operating Systems

---

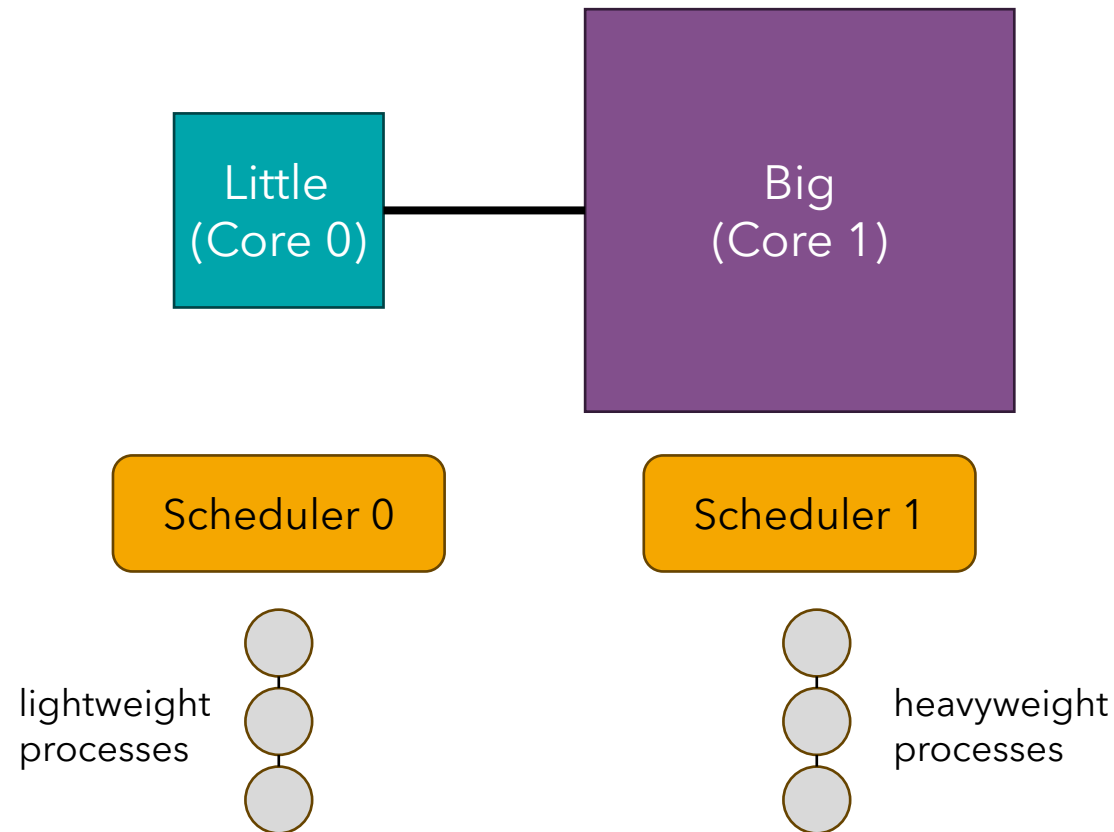
Project I - Scheduler

Minjae Kim

[jwya2149@dgist.ac.kr](mailto:jwya2149@dgist.ac.kr)

# Project #1

- Design & Implement a basic Big-Little CPU Scheduler
  - Modify the xv6 scheduler to a scheduler that manages workloads in a big-little multicore architecture



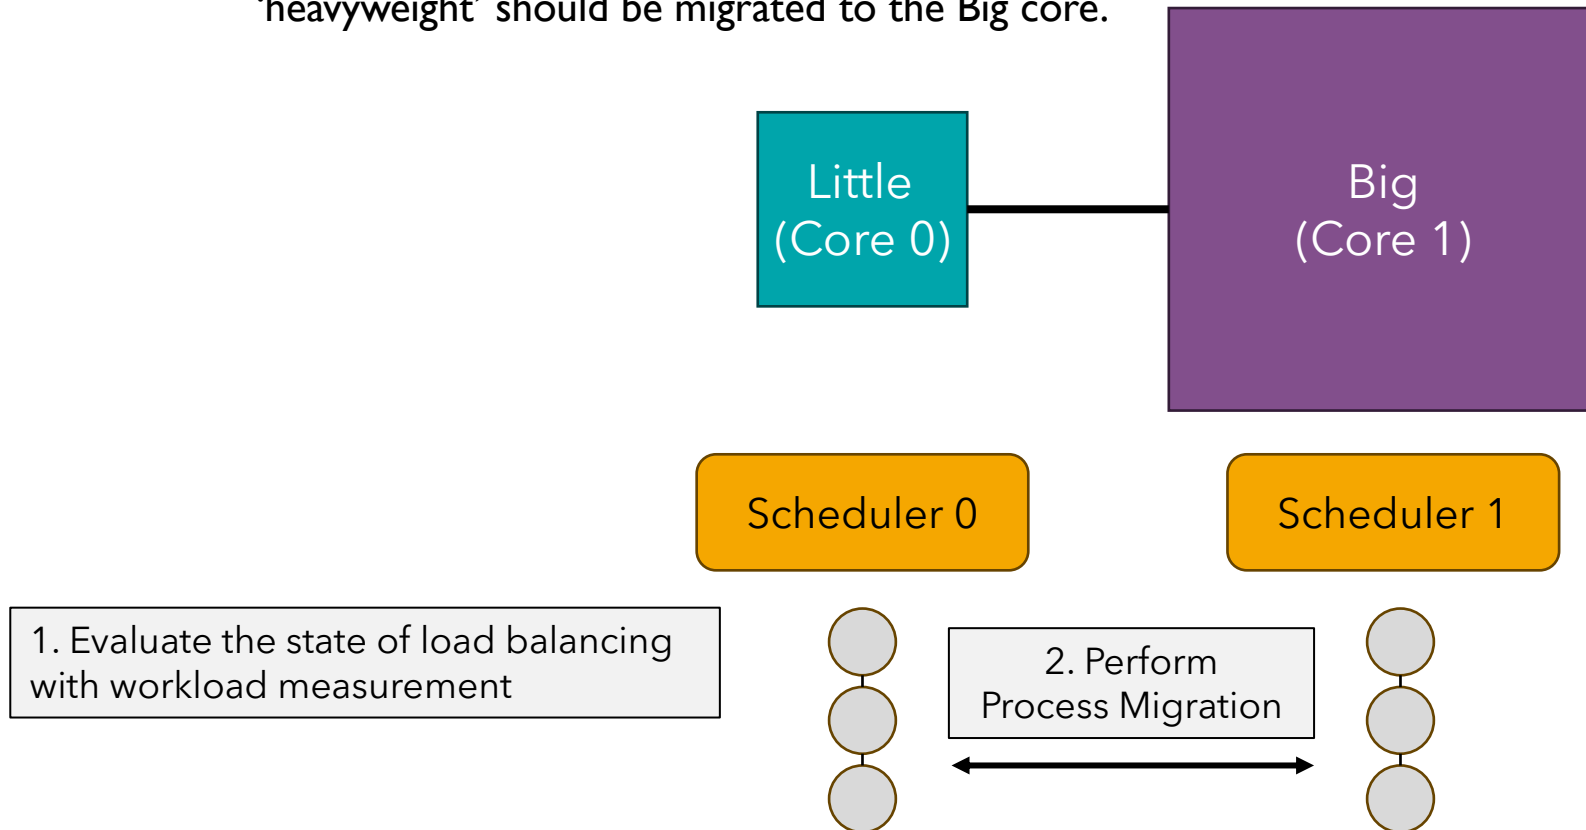


# Project #1

- Objectives of this project
  - Understand how context-switches are performed in xv6 code
  - Investigate the functioning of the big-little core architecture
  - Explore how to differentiate between heavy/light processes based on their resource requirements
- Where to look and write code:
  - `proc.c`, `proc.h` (+ etc)

# Project #1

- Implementation Approach in xv6
  - Assume there is one Little core and one Big core, with a per-CPU scheduler for each core
  - Newly created processes are assigned to the Little core
  - If the workload in the Little Scheduler exceeds a specified threshold, the processes that are classified as 'heavyweight' should be migrated to the Big core.



# Project #1

- Criteria for process migration? (i.e., Little core is busy)
  - # of RUNNABLE processes
  - Total CPU usage time
  - I/O wait Time
  - ...
- Criteria for heavyweight processes?
  - Execution time (i.e., CPU runtime)
  - Priority level
  - Memory usage
  - Historical performance data
  - ...

e.g., when the number of RUNNABLE processes in the Little core is over 10, select the process of the highest CPU runtime and migrate it to Big core.





# Tasks for Students

- Step 1
  - Assign newly created processes to CPU 0
  - Implement a function to enable process migration between CPUs
- Step 2
  - Develop a system load measurement feature
  - Implement a load balancing mechanism to identify heavy processes and migrate them to Big core



# Experimental Environments

- Benchmarks to be provided
  - Programs that generate multiple heavy and light processes
  - Measurements of scheduling metrics such as total time, turnaround time, response time, ...
  - Verification of whether heavy processes are executed on the Big Core
- Big/Little Core Simulation
  - Direct modification of CPU performance is not possible in xv6
  - The provided benchmarks will ensure that heavy workloads continuously run on CPU 0 to simulate Little Core behavior



# Evaluation

- Proper Functionality
  - Ensure the system operates without kernel panics
  - Verify that the system functions reasonably and efficiently
- Process Classification
  - Evaluate how effectively heavy and light processes are distinguished
- Report
  - Implementation Clarity: Assesses the clarity and thoroughness of your implementation description
  - Results Analysis: Analyze the results effectively, including insights on performance and behavior






# Project #1

- Deadline
  - ~ 2024.10.30 (Wed) 23:59
- Hand-in procedure
  - proj1\_202212345.patch (student number)
    - Run the following command and upload proj1\_202212345.patch
      - `git diff > proj1_202212345.patch`
    - Check the patch file with Notepad and confirm your modifications are in the patch file
  - Report
    - Submit a report (No page limit, but 2~3 pages are enough)
      - Free format (Korean/English)
      - Description of your implementation
      - Analysis of benchmark programs

# Project #1

- What is the default xv6 scheduler?
  - Round-robin scheduler
  - For each timer tick ( $\sim 10\text{ms}$ ), a timer interrupt occurs to incur a context switch
  - Don't need to change this scheduling policy



# Project #1

- Functions to create a process in xv6
  - proc.c: fork(), allocproc()

```
180 int
181 fork(void)
182 {
183     int i, pid;
184     struct proc *np;
185     struct proc *curproc = myproc();
186
187     // Allocate process.
188     if((np = allocproc()) == 0){
189         return -1;
190     }
191
```

...

```
216
217     np->state = RUNNABLE;
218
219     release(&ptable.lock);
220
221     return pid;
222 }
```

```
73 static struct proc*
74 allocproc(void)
75 {
76     struct proc *p;
77     char *sp;
78
79     acquire(&ptable.lock);
80
81     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
82         if(p->state == UNUSED)
83             goto found;
84
85     release(&ptable.lock);
86     return 0;
87
88 found:
89     p->state = EMBRYO;
90     p->pid = nextpid++;
91
92     release(&ptable.lock);
93
```

# Project #1

- Core function: `void scheduler()` in `proc.c`
- First *for* loop is looping forever
  - This function never returns
  - Find a new process to be scheduled, run it until it yields

**Infinite loop  
(1 tick)**

- Scan `ptable` to find RUNNABLE process

```
10 struct {  
11     struct spinlock lock;  
12     struct proc proc[NPROC];  
13 } ptable;
```

- Switch from scheduler to the process
  - After the process yields by timer interrupt, come back to `swtch()`

<default round-robin scheduler code>

```
322 void scheduler(void) {  
323     struct proc *p;  
324     struct cpu *c = mycpu();  
325     c->proc = 0;  
326  
327     for(;;){  
328         // enable interrupts on this processor.  
329         sti();  
330  
331         // Loop over process table looking for process to run.  
332         acquire(&ptable.lock);  
333         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){  
334             if(p->state != RUNNABLE)  
335                 continue;  
336  
337             // Switch to chosen process. It is the process's job  
338             // to release ptable.lock and then reacquire it  
339             // before jumping back to us.  
340             c->proc = p;  
341             switchvm(p);  
342             p->state = RUNNING;  
343  
344             swtch(&(c->scheduler), p->context);  
345             switchkvm();  
346  
347             // Process is done running for now.  
348             // It should have changed its p->state before coming back.  
349             c->proc = 0;  
350         }  
351         release(&ptable.lock);  
352     }  
353 }
```

# Project #1

- 10ms after the user process occupies the CPU, timer interrupt happens

in trap.c, trap() calls yield()

```
36 void
37 trap(struct trapframe *tf)
38 {
39     if(tf->trapno == T_SYSCALL){
40         if(myproc()->killed)
41             exit();
42         myproc()->tf = tf;
43         syscall();
44         if(myproc()->killed)
```

```
...
103 // Force process to give up CPU on clock tick.
104 // If interrupts were on while locks held, would need to check nlock.
105 if(myproc() && myproc()->state == RUNNING &&
106     tf->trapno == T_IRQ0+IRQ_TIMER)
107     yield();
```

```
384 // Give up the CPU for one scheduling round.
385 void
386 yield(void)
387 {
388     acquire(&ptable.lock); //DOC: yieldlock
389     myproc()->state = RUNNABLE;
390     sched();
391     release(&ptable.lock);
392 }
```

```
365 void
366 sched(void)
367 {
368     int intena;
369     struct proc *p = myproc();
370
371     if(!holding(&ptable.lock))
372         panic("sched ptable.lock");
373     if(mycpu()->ncli != 1)
374         panic("sched locks");
375     if(p->state == RUNNING)
376         panic("sched running");
377     if(readeflags() & FL_IF)
378         panic("sched interruptible");
379     intena = mycpu()->intena;
380     swtch(&p->context, mycpu()->scheduler);
381     mycpu()->intena = intena;
382 }
```

back to scheduler(), starts at line 345

# Project #1

- per-process structure: struct proc in proc.h


```
37 // Per-process state
38 struct proc {
39     uint sz;
40     pde_t* pgdir;
41     char *kstack;
42     enum procstate state;
43     int pid;
44     struct proc *parent;
45     struct trapframe *tf;
46     struct context *context;
47     void *chan;
48     int killed;
49     struct file *ofile[NOFILE];
50     struct inode *cwd;
51     char name[16];
52 };
```

- You can add member variables for per-process variable  
e.g.,) runtime, cpu\_affinity, ...

# Project #1

## Scheduler 0

```
322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     cprintf("Scheduler id: %d\n", c->apicid);
330
331     for(;;){
332         // Enable interrupts on this processor.
333         sti();
334
335         // Loop over process table looking for process to run.
336         acquire(&ptable.lock);
337         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
338             if(p->state != RUNNABLE)
339                 continue;
340
341             // Switch to chosen process. It is the process's job
342             // to release ptable.lock and then reacquire it
343             // before jumping back to us.
344             c->proc = p;
345             switchvm(p);
346             p->state = RUNNING;
347
348             swtch(&(c->scheduler), p->context);
349             switchkvm();
350
351             // Process is done running for now.
352             // It should have changed its p->state before coming back.
353             c->proc = 0;
354         }
355         release(&ptable.lock);
356     }
357 }
358 }
```



## Scheduler 1

```
322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     cprintf("Scheduler id: %d\n", c->apicid);
330
331     for(;;){
332         // Enable interrupts on this processor.
333         sti();
334
335         // Loop over process table looking for process to run.
336         acquire(&ptable.lock);
337         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
338             if(p->state != RUNNABLE)
339                 continue;
340
341             // Switch to chosen process. It is the process's job
342             // to release ptable.lock and then reacquire it
343             // before jumping back to us.
344             c->proc = p;
345             switchvm(p);
346             p->state = RUNNING;
347
348             swtch(&(c->scheduler), p->context);
349             switchkvm();
350
351             // Process is done running for now.
352             // It should have changed its p->state before coming back.
353             c->proc = 0;
354         }
355         release(&ptable.lock);
356     }
357 }
358 }
```



# Project #1

- References

- To understand more detail about xv6 scheduler, study Chapter 5 in the xv6 book (<https://pdos.csail.mit.edu/6.828/2018/xv6/book-rev11.pdf>)

- Build xv6 with 'CPUS = 2' flag to test easier (In Makefile)

```
219 ifndef CPUS
220 CPUS := 2
221 endif
```

mpmain() called twice?

```
50 // Common CPU setup code.
51 static void
52 mpmain(void)
53 {
54     cprintf("cpu%d: starting %d\n", cpuid(), cpuid());
55     idtinit(); // load idt register
56     xchg(&(mycpu()->started), 1); // tell startothers() we're up
57     scheduler(); // start running processes
58 }
```



Finally...

**Do NOT hesitate** to ask questions!