

# Operating Systems Practice

---

Project #2 – (Large) Inverted Page Table

Seonggyun Oh

([sungkyun123@dgist.ac.kr](mailto:sungkyun123@dgist.ac.kr))






# Project #2

- Implement a large inverted page table in the xv6
  - Modify the default paging scheme (multi-level paging) to large inverted page table
- Objectives of this project
  - Understand How multi-level paging are performed in xv6 code
  - Implement a inverted page table
- Where to look and write code:
  - kalloc.c, vm.c (+etc)

# Paging Structure in xv6: Multi-Level Paging

- 2-level paging
  - 1 page directory and 1 page table
  - Example: Write page to VA 0x9002000 (0b00001001000000000010000000000000)




Page directory  
address

0x0400 → PD	Index	PA	Valid
	0	0x10000	1
	1	-	0
	...	...	...
	36	-	0
	...	...	...

Page directory

# Paging Structure in xv6: Multi-Level Paging

- 2-level paging
  - 1 page directory and 1 page table
  - Example: Write page to VA 0x9002000 (0b00001001000000000010000000000000)



Page directory address

0x0400

PD	Index	PA	Valid
	0	0x10000	1
	1	-	0
	...	...	...
	36	-	0
	...	...	...
	...	...	...

PDE

Page directory

**Allocation!**

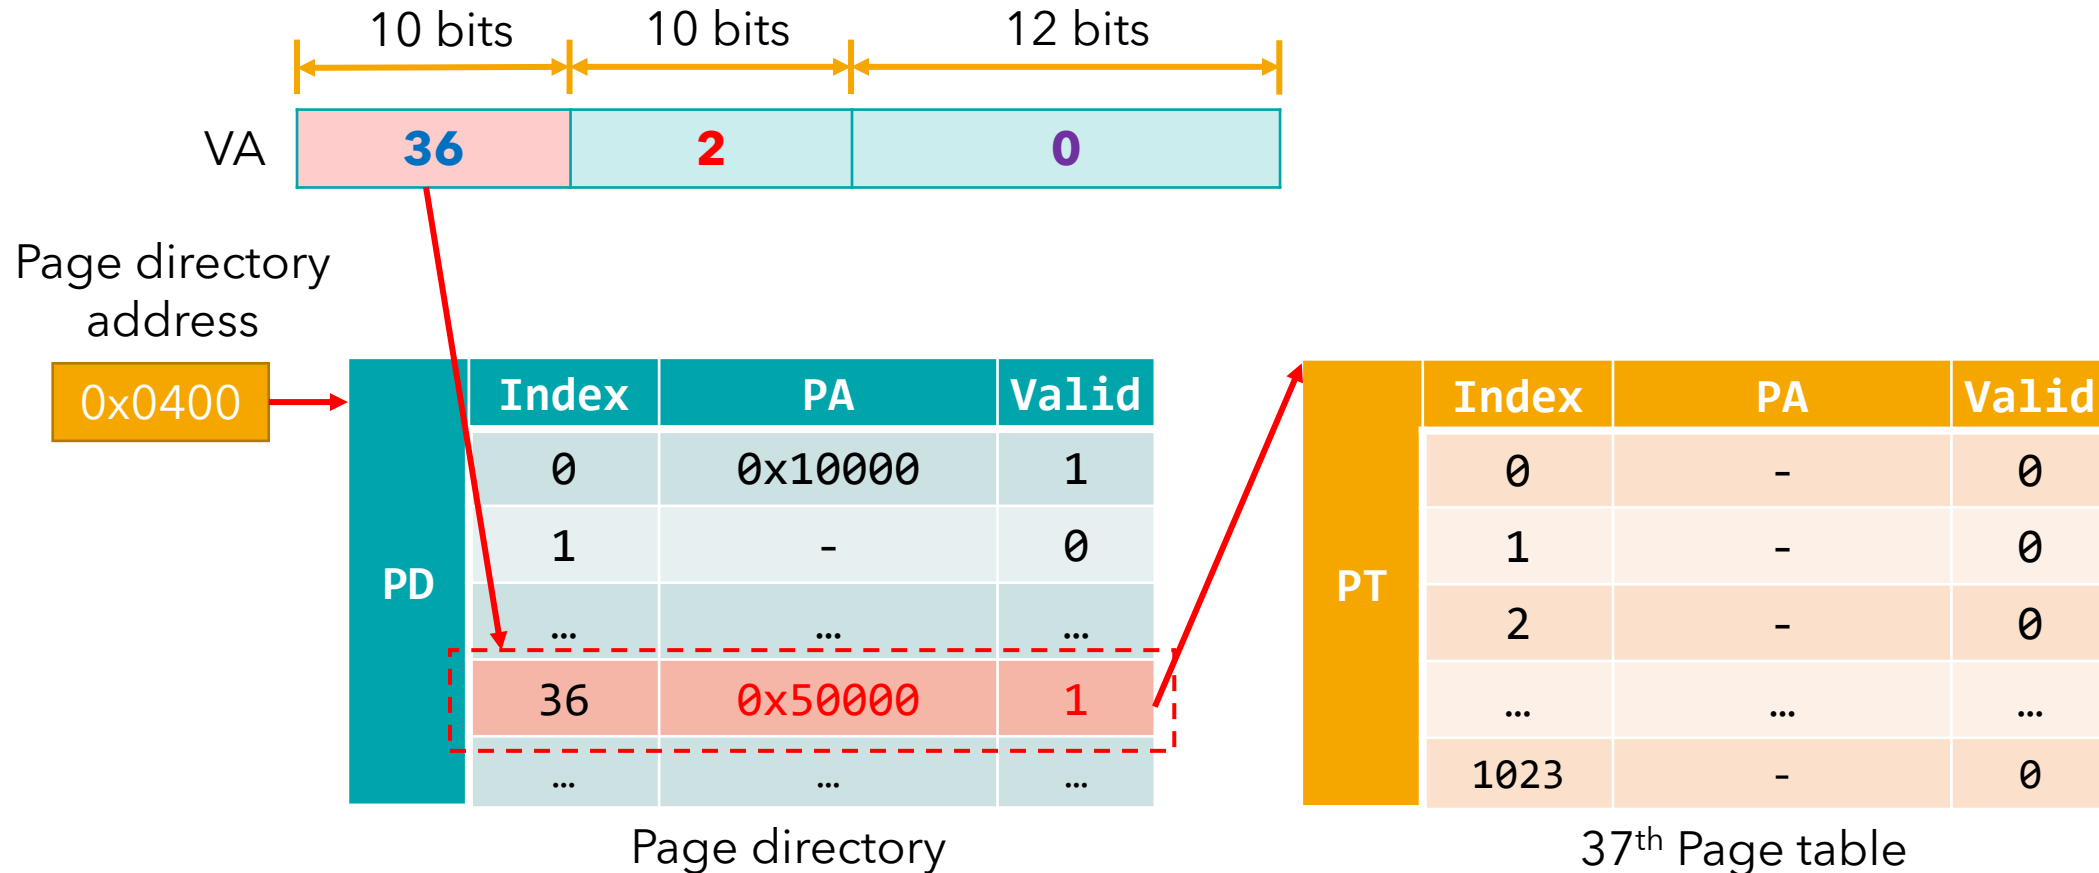
PT	Index	PA	Valid
	0	-	0
	1	-	0
	2	-	0
	...	...	...
	1023	-	0
	...	...	...

37<sup>th</sup> Page table




# Paging Structure in xv6: Multi-Level Paging

- 2-level paging
  - 1 page directory and 1 page table
  - Example: Write page to VA 0x9002000 (0b00001001000000000010000000000000)



# Paging Structure in xv6: Multi-Level Paging

- 2-level paging
  - 1 page directory and 1 page table
  - Example: Write page to VA 0x9002000 (0b0000|00|0000000000|000000000000)



Page directory  
address

0x0400

PD	Index	PA	Valid
	0	0x10000	1
	1	-	0
	...	...	...
	36	0x50000	1
	...	...	...

Page directory

PT	Index	PA	Valid
	0	-	0
	1	-	0
	2	-	0
	...	...	...
	1023	-	0

37<sup>th</sup> Page table


**Allocation!**

4K page

**PTE**

# Paging Structure in xv6: Multi-Level Paging

- 2-level paging
  - 1 page directory and 1 page table
  - Example: Write page to VA 0x9002000 (0b0000|00|0000000000|000000000000)



Page directory address

0x0400

PD	Index	PA	Valid
	0	0x10000	1
	1	-	0
	...	...	...
	36	0x50000	1
	...	...	...

Page directory

PT	Index	PA	Valid
	0	-	0
	1	-	0
	2	0x59000	1
	...	...	...
	1023	-	0

37<sup>th</sup> Page table

PTE

4K page



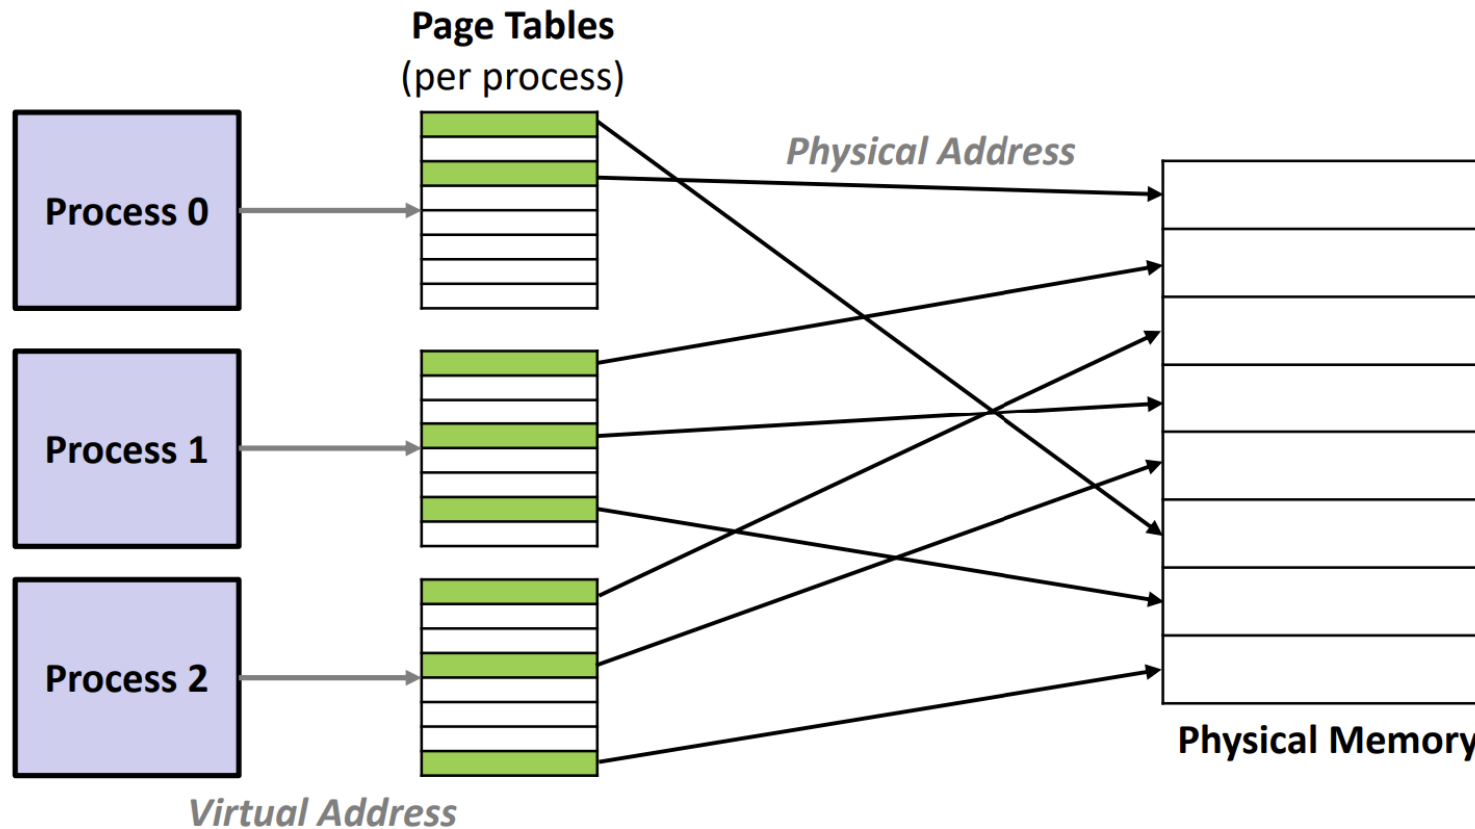
# Pros and Cons of Multi-Level Paging

- Pros
  - Memory saving compared to linear page table
    - Only allocates page-table space in proportion to the amount of address space the process uses
- Cons
  - Time-space trade-off
    - On a TLB miss, it requires multiple memory references
      - it gets worse as the number of levels increases
  - Complexity
    - It should manage more than one table
  - Memory requirement
    - As the number of processes increases, the memory requirement for paging is non-negligible



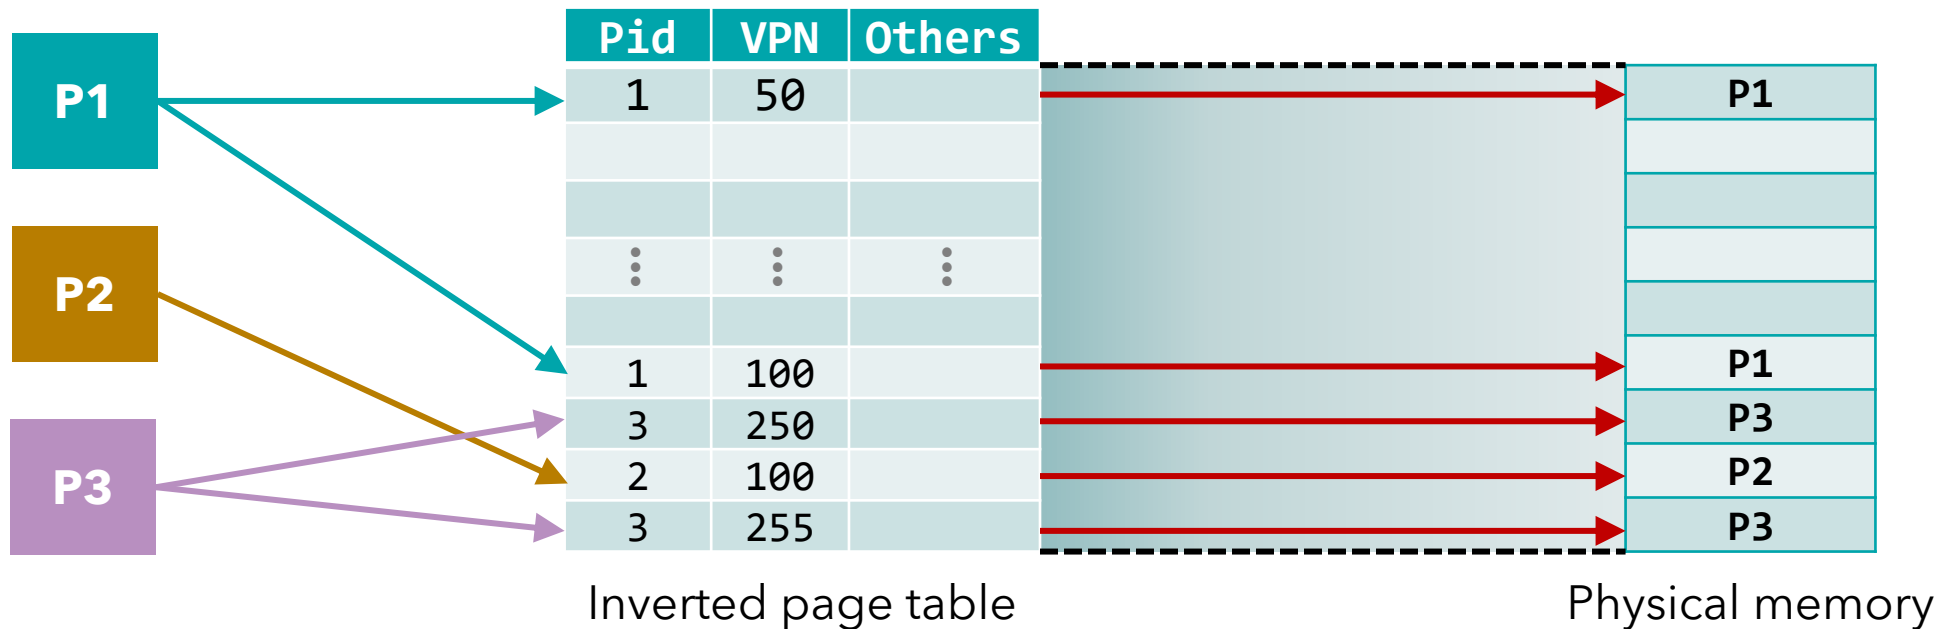
# Pros and Cons of Multi-Level Paging

- In the paging, all the processes maintain their own page table
  - As the number of processes increases, the wasted memory space increases



# Inverted Page Table: Overview

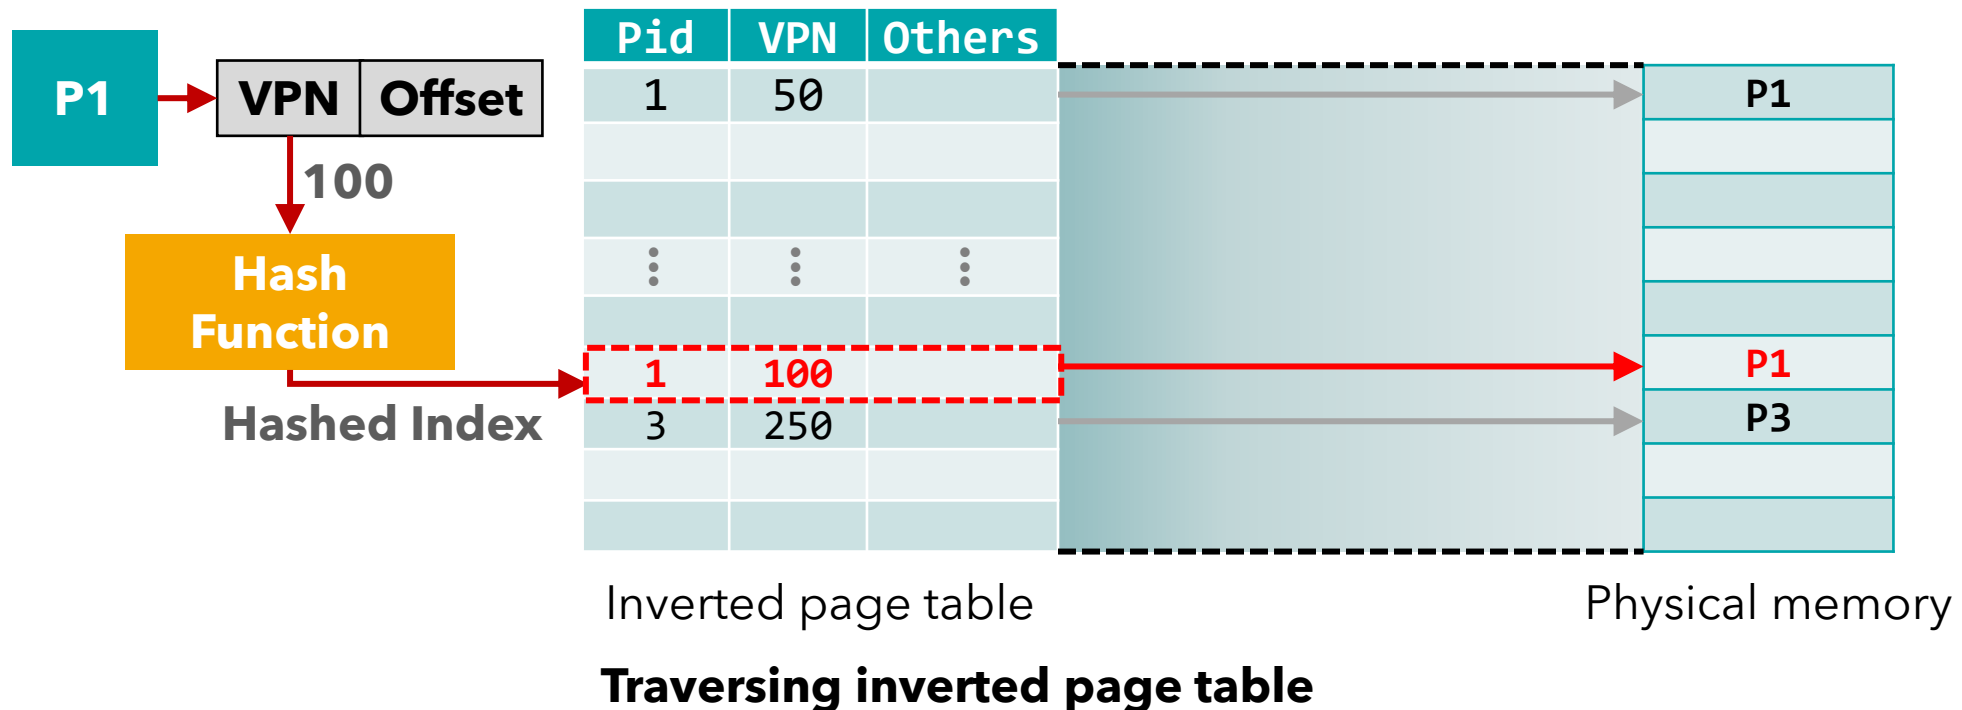
- Keeps a single page table that has a page table entry for each physical page of the system
  - Pros
    - Memory efficient – No need to maintain individual page tables for processes, thereby scaling with the size of physical memory



**# of PTEs = # of Physical pages**


# Inverted Page Table: Mapping process

- A virtual address is hashed to a specific PTE in the table
- The index of the PTE is equal to the page frame number of the page it maps




# Inverted Page Table: Mapping process

- A virtual address is hashed to a specific PTE in the table
  - **Hash collision** may occur (must handle!)
- The index of the PTE is equal to the page frame number of the page it maps




# Inverted Page Table: Mapping Process

- A virtual address is hashed to a specific PTE in the table
  - **Hash collision** may occur (must handle!)
- The index of the PTE is equal to the page frame number of the page it maps



# Inverted Page Table: Summary


- Pros
  - Lower memory consumption
- Cons
  - Long tail latency by hash collision... **How to improve?**






# Large Inverted Page Table: Overview

- Main concept: Increase the number of entries in the inverted page table to reduce hash collision!
- Then, the index of the inverted page table will no longer be the physical page frame number




# Large Inverted Page Table: Overview

- Difference from the original inverted page table
  - New column “**PPN**”: Record the physical page number
  - **Freelist**: Linked list composed of pointers to free physical memory pages



- **Remove the page at the tail of the freelist, freeing it for use**
- Using hash function and linear probing, find the proper page table entry
- Record the PID, VPN and PPN of freed page



- Remove the page at the tail of the freelist, freeing it for use



# VPN

# Offset

# Hash Function

900

## Hashed Index

## Large Inverted page table

[illegible]

## Freelist

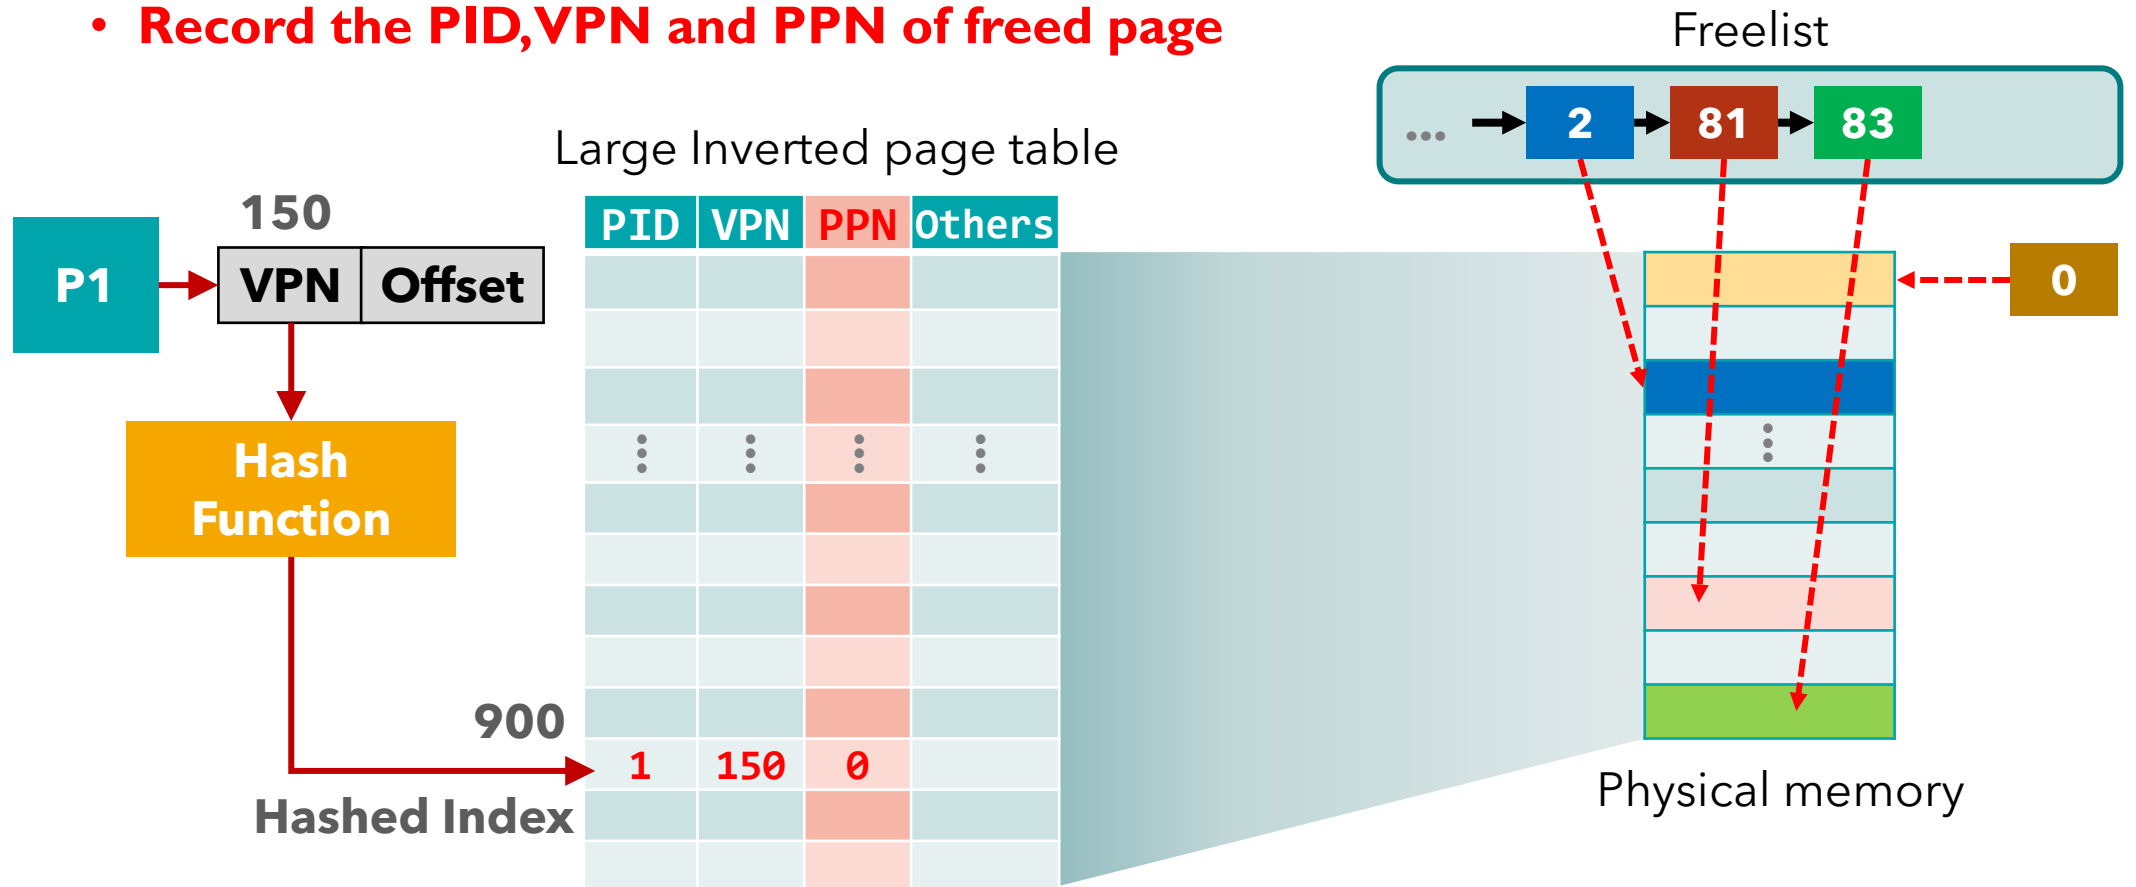
... → **2** → **81** → **83**

0

## Physical memory


# Large Inverted Page Table: Mapping Process

- Memory page allocation
  - Remove the page at the tail of the freelist, freeing it for use
  - Using hash function and linear probing, find the proper page table entry
  - **Record the PID,VPN and PPN of freed page**



# Large Inverted Page Table: Mapping Process


- Memory page deallocation (free)
  - Find the PPN for given PID and VPN
  - Erase the content of the corresponding page (memset)
  - Initialize the page table entry
  - Insert the free page into Freelist






# Large Inverted Page Table: Mapping Process

- Memory page deallocation (free)
  - Find the PPN for given PID and VPN
  - **Erase the content of the corresponding page (memset)**
  - Initialize the page table entry
  - Insert the free page into Freelist




# Large Inverted Page Table: Mapping Process

- Memory page deallocation (free)
  - Find the PPN for given PID and VPN
  - Erase the content of the corresponding page (memset)
  - **Initialize the page table entry**
  - Insert the free page into Freelist



- **Insert the free page into Freelist**



# Project #2: Implementing Large Inverted Page Table

- Virtual Memory for supporting inverted page table
  - Allocating and Freeing
    - When `kalloc()` is called, given pid and virtual address, you must allocate the proper page
    - When `kfree()` is called, you must find the corresponding PTE and free that page
    - If hash collision occur, you must handle that by linear probing
    - Where to implement: **`kalloc()` and `kfree()`** function in `kalloc.c` and you can add functions in that code
  - Traverse
    - A virtual address is hashed to a specific PTE in the table
    - If hash collision occurs, perform linear probing
    - Where to implement: **`ittraverse()`** function in `vm.c`
  - Deallocate
    - When the process is terminated, free allocated pages and clear the corresponding entry
    - Where to implement: **`deallocvm()`** function in `vm.c`

# I. Allocating and Freeing Page

- There are four structures: the number of entries for each structure is 40000
  - PID[i]: Stores a PID whose process requests allocation of page i
  - VPN[i]: Stores a VPN that corresponds to page i
  - PPN[i]: Stores a PPN of physical page taken from Freelist
  - PTE\_XV6[i]: This structure is for compatibility of xv6, and this structure stores the physical address with permission, (same as PTE of multi-level paging in original xv6)
- Allocation (kalloc(pid, v))
  - Get free page from freelist
  - Find the empty entry using hash function and linear probing
  - Record the PID, VPN, PPN into the page table entry
  - Return (char\*)r; (r is the P2V(physical address))
  - Edge case: if  $v == -1$ , set vpn to P2V(physical address)
    - Here, v is the virtual address

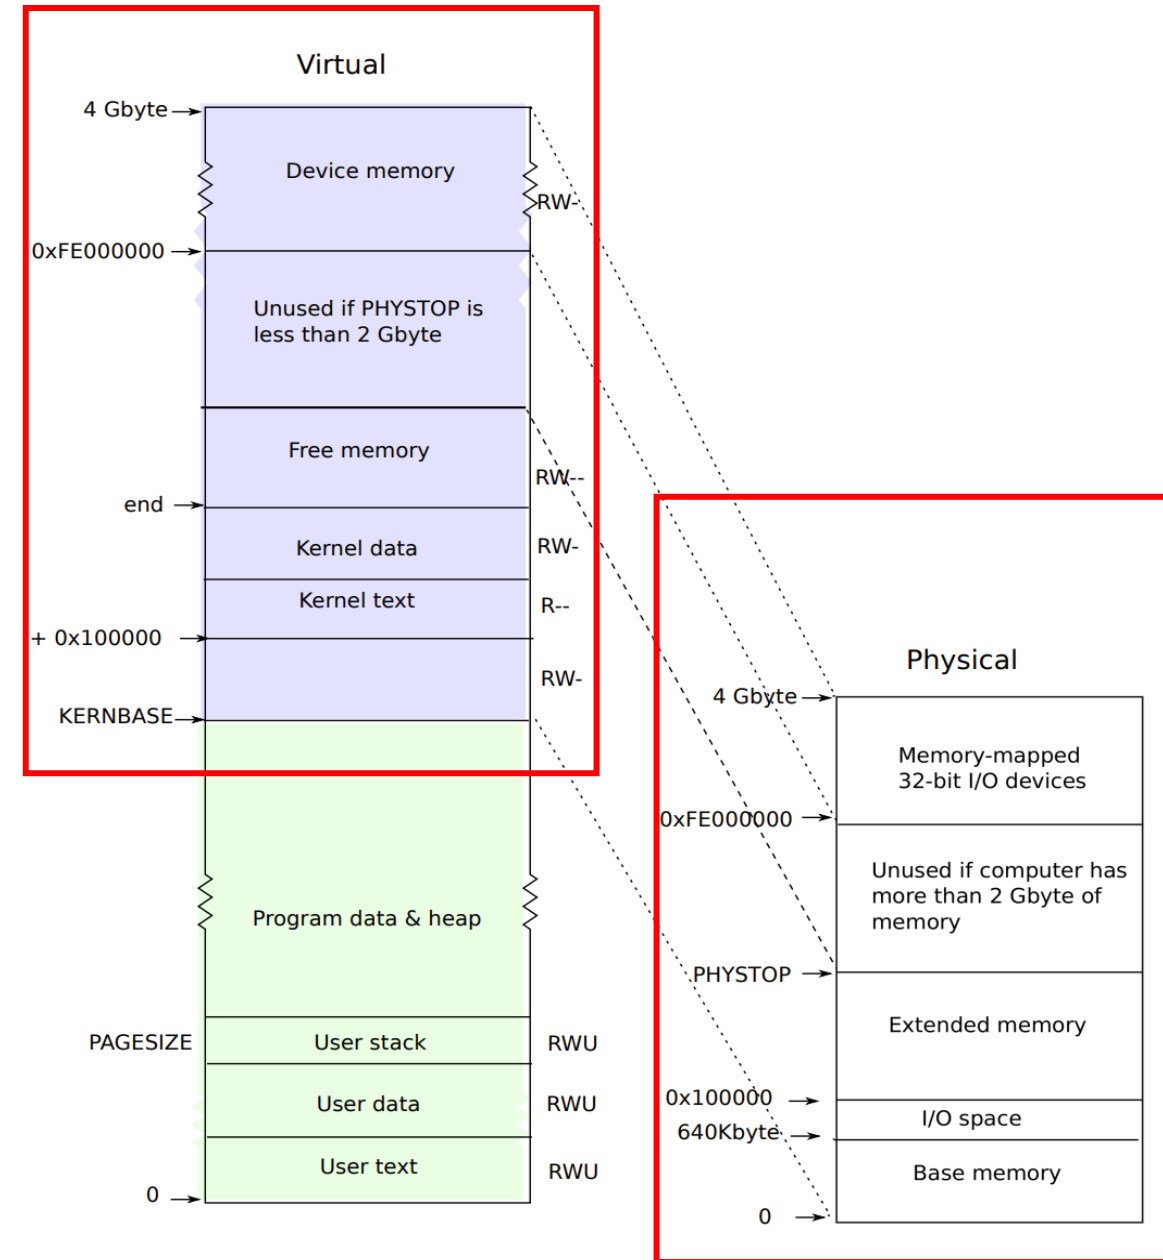
# I. Allocating and Freeing Page

- There are four structures: the number of entries for each structure is 40000
  - PID[i]: Stores a PID whose process requests allocation of page i
  - VPN[i]: Stores a VPN that corresponds to page i
  - PPN[i]: Stores a PPN of physical page taken from Freelist
  - PTE\_XV6[i]: This structure is for compatibility of xv6, and this structure stores the physical address with permission, (same as PTE of multi-level paging in original xv6)
- Deallocation (kfree)
  - Find the target entry using hash function and linear probing
  - Initialize the PID, VPN, PPN into the page table entry
  - Insert erased page into Freelist



## 2.Traverse

- You must consider Kernel address space
  - Kernel address space is directly mapped to physical memory space by subtracting VA from KERNBASE (0x8000000)
- **ittraverse()** function in vm.c
  - If  $VA \geq KERNBASE$ :
    - Return  $\&PTE\_KERN[V2P(VA)/PGSIZE]$
  - Else?
    - Return  $\&PTE\_XV6[idx]$
    - Idx: corresponding index for given PID and VPN



**Virtual and Physical memory space in xv6**

## 2.Traverse

- **ittraverse()** function in vm.c
  - If  $VA > KERNBASE$ :
    - Return: `&PTE_KERN[V2P(VA)]`
  - Else?
    - **Very easy! You just traverse the inverted page table for given pid and virtual address (VR)**
      - Make a function for general purpose (`searchit()`)
    - Return: `&PTE_XV6[idx]`
- Why use `PTE_XV6`?
  - In original xv6 code that support multi-level paging, the `ittraverse()` return page table entry itself, which can contain physical address and permission.
  - To minimize the modification of xv6, our implementation follows same format of output

```
ittraverse()
{
    if VA > KERNBASE:
        return &PTE_KERN[V2P(VA)/PGSIZE];

    else:
        idx = searchit(va, pid);
        return &PTE_XV6[idx];
}
```

**Pseudocode of ittraverse()**

### 3. Deallocating User Virtual Memory

- When you deallocate the user process's virtual memory space, you must find how the process requests allocation for kernel
  - For example, When process A has 10 allocated pages for kernel, you should find them in the inverted page table
  - **Deallocvm() in vm.c**
    - For a given range, deallocate the previous-allocated pages within the range
    - How?
      - **By your own idea!**
        - For example, traverse whole inverted page table and when the condition for deallocation, call kfree()

# Testing

- To evaluate your implementation, you should perform **usertests** and **hashtests**
- You should capture the result of usertests
  - Just type “usertests” in xv6 console
- In hashtests, linear probing occurs. You must capture the linear probing case
  - Print this information
    - Example (PID and VA for requests)
      - [Hash collision for idx: 2] PID: 0, VA: 0x100000, VA is different
      - [Hash collision for idx: 3] PID: 0, VA: 0x100000, PID is different
      - [Hash collision for idx: 4] PID: 0, VA: 0x100000, PID is different
      - [Completion idx: 5] PID: 0, VA: 0x100000
  - Change the MAXENTRY from 50000 to 70000 (in param.h)
    - Perform hashtests and analyze the impact of inverted page table size

# Summary

- You should modify **allocation, deallocation, traversal code** to support **Inverted page table** in context of user's process
  - If needed, you can add functions
    - `kalloc()`, `kfree()`, `ittraverse()`, `deallocvm()`
- To evaluate your implementation, you should perform **usertests** and **hashtests()**
- Fill in the code wherever it contains *//TODO*

# Project #2 – Inverted Page Table

- Deadline
  - ~ 2024.11.25 (Mon) 23:59
- Project repo. Generation
  - mkdir os-pj2; cd os-pj2; git clone <https://github.com/dgist-datalab/xv6>; cd xv6; git checkout inverted
- Hand-in procedure
  - p2\_201812345.patch
  - Run the following command and upload p2\_201812345.patch
    - git diff > p2\_201812345.patch
  - Check the patch file with Notepad and confirm your modifications are in the patch file
  - Report (Important!)
    - Submit an 1~3 pages report
      - Free format (Korean/English)
      - Description of your implementation in detail (kalloc, kfree, ittraverse, deallocvm)
      - Explain your test code and answer the following two tests and explain why the results are as follows
      - Insert test code result image





Finally ...

**Do NOT hesitate** to ask questions!

Mini Project #1, #2   Juhyung Park   *arter97@dgist.ac.kr*

Project #1   Minjae Kim   *jwya2149@dgist.ac.kr*

**Project #2   Seonggyun Oh   *sungkyun123@dgist.ac.kr***

Project #3   Jeeyun Kim   *kimgyun@dgist.ac.kr*